

Project Report

on

Advanced Phishing Detection System Using Hybrid AI Models

Submitted in partial fulfillment of requirements for the award of the degree

Bachelor of Technology

In

Artificial Intelligence & Data Science

To

IKG Punjab Technical University, Jalandhar

Submitted By:

Name: Suhail Kataria

Roll no.: 2417163

Semester: 6th

Batch: 2023-2027

Under the guidance of

Ms. Nidhi Bhatla

Assistant Professor



**CGC COLLEGE
OF ENGINEERING**

Building Careers. Transforming Lives.

Department of Computer Science & Engineering

CGC – College of Engineering, Landran

Mohali, Punjab – 140307

April 2026

CANDIDATE’S DECLARATION

I hereby declare that the work which is being presented in the 6th semester, entitled **“Advanced Phishing Detection System Using Hybrid AI Models”**, in partial fulfillment of the requirements for the award of the degree of Bachelor of Technology in Computer Science & Engineering and submitted to CGC College of Engineering, Landran, Mohali is an original piece of work carried out by me during the period from January 2026 to May 2026 under the supervision of **Ms. Nidhi Bhatla**, department of CSE, CGC College of Engineering Landran Mohali Punjab India.

The matter embodied in this File has not been submitted by me for the award of any other degree of any other University/Institute.

Suhail Kataria

The File entitled **Advanced Phishing Detection System Using Hybrid AI Models** submitted by **Suhail Kataria** to Department of CSE for the award of the degree of Bachelor of Technology is a Bonafide record of the project work carried out by him/her under my supervision and guidance. This is to certify that the above statement made by the candidate is correct to the best of my knowledge.

Ms. Nidhi Bhatla
Assistant Professor
Department of CSE
CGC College of Engineering
Landran, Mohali, Punjab, India

Certificate

I hereby declare that the 6th Semester Training Report entitled “**Advanced Phishing Detection System Using Hybrid AI Models**” is an authentic record of my own work as requirements of 6th sem academic during the period from January 2026 to May 2026 for the award of degree of B. Tech. (**Artificial Intelligence & Data Science**), CGC-College of Engineering, Landran, Mohali under the guidance of **Ms. Nidhi Bhatla**

Suhail Kataria

2417163

Date: _____

Certified that the above statement made by the student is correct to the best of our knowledge and belief.

Signatures

Examined by:

1. 2. 3. 4.

Head of Department

ACKNOWLEDGMENT

With profound respect, sincere gratitude is expressed to Director-Principal **Dr. Anuj Kumar Gupta**, CGC College of Engineering, Landran, for providing the opportunity and institutional support to complete 6th Semester Project and complete the present work.

Heartfelt thanks are extended to **Dr. Sushil Kamboj**, Head, Department of Computer Science & Engineering, CGC College of Engineering, Landran (Mohali), for academic guidance, continuous encouragement, and valuable inputs throughout the training period.

Appreciation is also conveyed to the faculty members of the Computer Science & Engineering Department, CGC COE, Landran, for their constructive feedback, technical insights, and consistent support that strengthened the quality of this synopsis.

The assistance of **Ms. Nidhi Bhatla**, Supervisor for experimentation, is gratefully acknowledged for their mentorship, practical guidance, and timely facilitation of experimental work.

Finally, indebtedness is expressed to all individuals who have directly or indirectly contributed to the successful completion of this 6th Semester Training.

Suhail Kataria

Table of Content

Candidate Declaration.....	ii
Certificate.....	iii
Acknowledgement.....	iv
Table of Content.....	v-vii
List of Figures.....	viii
List of Tables.....	ix
List of Abbreviations.....	x
Abstract.....	xi
Chapter 1 : Introduction.....	1-4
1.1. Overview of the Project.....	2-3
1.2 Motivation.....	3-4
1.3 Organization of the Report.....	4
Chapter 2: Literature Survey.....	5-11
2.1 Traditional Machine Learning Approaches.....	7
2.2 Deep Learning-Based Approaches.....	7-8
2.3 URL-Based Phishing Detection.....	8-9
2.4 Hybrid Phishing Detection Systems.....	9
2.5 Challenges and Research Gaps.....	9-10
2.6 Summary of Literature Review.....	10-11
Chapter 3: Problem Statement.....	12-13
Chapter 4: Rationale/Need for the Project.....	14-17
4.1 Increasing Threat of Phishing Attacks.....	15
4.2 Limitations of Existing Systems.....	16

4.3	Need for Hybrid Detection Approach.....	16
4.4	Importance of Explainable Systems.....	16-17
4.4	Practical Relevance and Applications.....	17
Chapter 5: Objectives.....		18-20
5.1	Primary Objectives.....	18-19
5.2	Technical Objectives.....	19
5.3	System Objectives.....	19-20
5.4	User-Oriented Objectives.....	20
5.5	Overall Objective Summary.....	20
Chapter 6: Feasibility Study.....		21-25
6.1	Technical Feasibility.....	21-22
6.2	Economic Feasibility.....	22-23
6.3	Operational Feasibility.....	23
6.4	Implementation Feasibility.....	24
6.5	Scalability and Maintainability Feasibility.....	24-25
6.6	Summary of Feasibility.....	25
Chapter 7: Methodology.....		26-31
7.1	System Overview.....	26
7.2	Data Collection and Preprocessing.....	27
7.3	Email Content Analysis.....	27-28
7.4	URL Feature Analysis.....	28
7.5	Feature Engineering.....	28-29
7.6	Joint Detection Mechanism.....	29
7.7	Threshold Optimization.....	30-31
7.8	System Workflow.....	31
7.9	Summary of Methodology.....	31

Chapter 8: Project Implementation Plan/Timeline.....	33-35
8.1 Development Phases.....	33-34
8.2 Project Timeline Table.....	34
8.3 Summary of Timeline.....	35
Chapter 9: Results and Discussions.....	36-43
9.1 Evaluation Metrics.....	36
9.2 Email Model Performance.....	37
9.3 URL Model Performance.....	38-39
9.4 Joint Model Performance.....	39-40
9.5 Optimized Joint Model Performance.....	40
9.6 Confusion Matrix Analysis.....	41
9.7 Error Analysis.....	42
9.8 Discussion.....	43
9.9 Summary of Results.....	43
Chapter 10: Conclusion.....	44-45
Chapter 11: Conclusion and Future Scope.....	46-48
11.1 Advanced Model Improvements.....	46
11.2 Real-Time Threat Detection.....	46-47
11.3 Improved Feature Engineering.....	47
11.4 System Integration and Deployment.....	47
11.5 User Interface Enhancements.....	48
11.6 Expansion to Other Attack Types.....	48
11.7 Summary of Future Scope.....	48
References.....	49-50

LIST OF FIGURES

Figure 2.1	Comparison of Phishing Detection Approaches.....	11
Figure 6.1	Feasibility Summary.....	25
Figure 7.1	Workflow architecture of Joint detection Engine.....	30
Figure 7.2	Workflow diagram of final system.....	32
Figure 9.1	Workspace Input Window.....	39
Figure 9.2	Workspace Output Window.....	40
Figure 9.3	Accuracy comparison graph.....	41

LIST OF TABLES

Table 8.1	Project Timeline Table.....	34
Table 9.1	Email Model Performance.....	37
Table 9.2	URL Model Performance.....	38
Table 9.3	Joint Model Performance.....	39
Table 9.4	Optimized Joint Model Performance.....	40

LIST OF ABBREVIATIONS

AI	Artificial Intelligence
ML	Machine Learning
DL	Deep Learning
NLP	Natural Language Processing
URL	Uniform Resource Locator
HTTP	Hypertext Transfer Protocol
HTTPS	Hypertext Transfer Protocol Secure
IP	Internet Protocol
CNN	Convolutional Neural Network
TF-IDF	Term Frequency–Inverse Document Frequency
ROC	Receiver Operating Characteristic
BERT	Bidirectional Encoder Representations from Transformers
F1-Score	Harmonic Mean of Precision and Recall

ABSTRACT

Phishing remains one of the most widespread and damaging cyber threats, exploiting deceptive emails and malicious URLs to obtain sensitive information such as login credentials, financial records, and personal data. Many existing phishing detection systems focus on either email content analysis or URL-based detection in isolation, which limits their effectiveness in realistic environments where attackers often combine multiple deceptive techniques within a single campaign.

This project presents an advanced phishing detection framework that unifies email content analysis and URL intelligence through a hybrid artificial intelligence approach. The system employs transformer-based models, particularly DistilBERT, to capture contextual and semantic patterns in email text, while also incorporating engineered features derived from URL structure, domain properties, and behavioral indicators. Detection performance is further strengthened through a joint decision mechanism that combines outputs from the email and URL models using rule-based optimization and threshold tuning.

The proposed framework is evaluated on several publicly available datasets, including the Enron dataset, the Nazario phishing corpus, the SpamAssassin dataset, PhishTank URLs, and the Top-1M benign URL dataset. The experimental results indicate that although traditional models often struggle under cross-domain conditions, the hybrid approach offers stronger robustness and improved generalization. The optimized joint detection system achieves more than 90% accuracy on an expanded evaluation dataset, suggesting strong potential for real-world deployment.

In addition, the system supports multiple operational modes, including high-recall detection for security operations environments and high-confidence alerting for more critical settings. The inclusion of explainable outputs further improves its practical value for cybersecurity analysts by making model decisions easier to interpret. Overall, this work contributes a scalable, effective, and practically relevant solution for modern phishing detection by integrating contextual language with structural URL analysis.

CHAPTER – 1

Introduction

Phishing has emerged as one of the most pervasive and harmful forms of cyber attack in today's digital environment. It is a deceptive practice in which attackers use fraudulent emails, messages, or websites to manipulate users into disclosing sensitive information, including login credentials, financial records, and personal data. As internet usage and digital communication continue to expand, phishing attacks have also evolved in complexity and sophistication, making them increasingly difficult to detect through traditional security approaches.

Conventional phishing detection methods generally rely on either email content analysis or URL-based analysis as separate strategies. Email-focused systems typically attempt to identify suspicious keywords, linguistic patterns, and sender-related anomalies, whereas URL-based approaches examine domain characteristics, structural patterns, and reputation signals. Although both methods can be effective in limited settings, they often prove insufficient in realistic threat environments where attackers combine multiple evasion techniques within a single campaign. As a result, systems that depend on a single detection strategy frequently lack robustness and perform poorly when exposed to diverse real-world phishing patterns.

A further challenge in phishing detection lies in the variability of datasets and the constantly changing nature of attack behavior. Models that perform well on one dataset often fail to maintain the same level of performance when tested on another, indicating limited generalization and a tendency to overfit source-specific patterns. This issue highlights the need for a more comprehensive and adaptable detection framework capable of handling variation across datasets, writing styles, and attack strategies.

To address these limitations, this project proposes a hybrid phishing detection system that integrates both email content analysis and URL intelligence within a unified framework. The system employs advanced machine learning techniques, including transformer-based

models such as DistilBERT, to capture contextual and semantic information from email text. In addition, engineered features derived from URLs and structural characteristics of emails are incorporated to strengthen detection capability and improve overall classification performance.

To further enhance effectiveness, a joint detection mechanism is introduced to combine the outputs of the email and URL models. This mechanism applies rule-based optimization and threshold tuning in order to improve decision quality and overall system reliability. By integrating multiple signals and complementary detection strategies, the proposed system is designed to achieve greater robustness and stronger generalization across heterogeneous datasets.

The system is evaluated using several publicly available datasets, including the Enron dataset, the Nazario phishing corpus, the SpamAssassin dataset, PhishTank URLs, and the Top-1M benign URL dataset. Experimental findings demonstrate clear improvements in detection performance, particularly under cross-source evaluation settings where traditional models often struggle. These results indicate that the hybrid approach is better suited to handling the variability and complexity of modern phishing attacks.

This project aims to provide a scalable and practical solution for phishing detection that can be integrated into real-world cybersecurity environments. By combining contextual understanding, structural analysis, and joint decision-making, the proposed system offers an effective approach for detecting phishing threats in contemporary digital ecosystems.

1.1 Overview of the Project

This project is centered on the design and implementation of a hybrid phishing detection system that combines email content analysis with URL-based intelligence. The proposed system is intended to address the limitations of traditional single-model approaches by bringing multiple detection strategies together within a unified framework.

The email analysis component employs transformer-based models, specifically DistilBERT, to capture contextual patterns and semantic relationships from the subject

and body of an email. This enables the system to identify phishing attempts even when attackers avoid obvious keywords and instead rely on subtle, persuasive, or context-aware language.

At the same time, the URL analysis component examines structural and domain-level characteristics of links embedded within emails. Features such as domain consistency, URL length, suspicious lexical patterns, and brand-domain mismatches are extracted and analyzed to detect indicators of malicious intent.

To improve overall performance, a joint detection mechanism is incorporated to combine the outputs of the email and URL models. This mechanism uses rule-based optimization and threshold-based decision strategies to support more reliable detection across different scenarios. In addition, the system is designed to operate under multiple modes, including high-recall detection for continuous security monitoring and high-confidence alerting for critical operational settings.

1.2 Motivation

The primary motivation for this project arises from the increasing sophistication and complexity of phishing attacks. Modern phishing campaigns rarely rely on a single deceptive element; instead, attackers often combine misleading email content with carefully crafted malicious URLs in order to evade conventional detection systems. In such an environment, relying on a single detection strategy is no longer sufficient.

Another important motivation is the limited generalization ability of many existing machine learning models. Although some models achieve strong performance on specific benchmark datasets, they often fail when evaluated on previously unseen data from different sources. This pattern suggests overfitting to dataset-specific characteristics and underscores the need for a more robust and adaptable detection framework.

In addition, there is a growing demand for practical and scalable cybersecurity solutions that can be deployed effectively in real-world environments. Security analysts require systems that not only produce accurate predictions, but also provide explainable outputs

and support flexible operating modes tailored to different security requirements.

This project seeks to address these concerns by developing a hybrid phishing detection system that improves classification accuracy, enhances robustness across domains, and offers practical value for real-world cybersecurity applications.

1.3 Organization of the Report

This report is organized into multiple chapters to provide a structured and comprehensive presentation of the proposed phishing detection system.

Chapter 1: Introduction presents the background of the study, the problem context, and the overall objectives of the project.

Chapter 2: Literature Survey reviews existing research, methods, and technologies related to phishing detection.

Chapter 3: Problem Statement defines the key challenges and issues that the project seeks to address.

Chapter 4: Rationale/Need for the Project explains the significance, relevance, and practical necessity of the proposed system.

Chapter 5: Methodology describes the overall design, architecture, and operational workflow of the system.

Chapter 6: Implementation discusses the tools, technologies, and development process used to build the project.

Chapter 7: Results and Discussion presents the evaluation metrics, experimental findings, and performance analysis of the proposed approach.

Chapter 8: Future Scope outlines possible improvements, extensions, and directions for further work.

Chapter 9: References lists the sources, datasets, and research materials used throughout the project.

CHAPTER – 2

LITERATURE SURVEY

Phishing detection has remained an important area of research in cybersecurity because of the growing frequency, scale, and sophistication of cyber attacks. Over time, a wide range of approaches has been proposed, beginning with traditional machine learning techniques and gradually evolving toward deep learning and hybrid detection frameworks designed to address more complex threat patterns.

Early research in phishing and spam detection relied heavily on classical machine learning algorithms such as Naïve Bayes and Logistic Regression. Metsis et al. [1], for example, investigated spam filtering using Naïve Bayes and demonstrated its effectiveness in text classification tasks. These early methods were attractive because of their computational efficiency and simplicity of implementation. However, they were generally limited in their ability to capture deeper contextual relationships within textual data, which reduced their effectiveness against more sophisticated phishing messages.

As machine learning methods advanced, researchers began to incorporate richer feature engineering strategies and more robust classification techniques. Studies such as “*AI Based Phishing Detection Techniques: A Comparative Analysis of Model Performance*” [2] highlighted the value of combining multiple feature types, including textual, lexical, and structural attributes, to improve detection accuracy. In a similar direction, hybrid machine learning approaches [3][4] demonstrated that integrating different models or feature sets could often outperform single-model systems by capturing a wider range of phishing indicators.

More recently, phishing detection research has increasingly shifted toward deep learning approaches. Devlin et al. [6] introduced BERT, a transformer-based language model capable of capturing contextual relationships and semantic meaning in text more effectively than traditional feature-based methods. Building on this work, DistilBERT [7] offers a lighter and more computationally efficient variant that is better suited for practical deployment while retaining much of the representational strength of larger

transformer models. Other deep learning architectures, including Convolutional Neural Networks (CNNs) and Long Short-Term Memory (LSTM) networks, have also been applied to phishing detection tasks [7], showing improved capability in identifying complex and less obvious attack patterns.

In addition to email content analysis, URL analysis has emerged as a critical component of phishing detection. Research such as “*Hybrid AI Models for Phishing URL Detection*” [8], along with studies based on datasets like PhishTank and Top-1M URLs, has demonstrated the importance of examining domain structure, lexical characteristics, and behavioral indicators within URLs. These methods are particularly valuable for identifying malicious links embedded in emails, which often serve as the main mechanism for redirecting users to fraudulent websites.

Despite these advances, one of the most persistent limitations reported in the literature is the problem of cross-domain generalization. Models trained on a particular dataset often experience a noticeable decline in performance when evaluated on previously unseen data from different sources. This issue has been emphasized in several surveys and comparative studies [2][5], which suggest that many models tend to overfit dataset-specific patterns rather than learning broadly transferable phishing characteristics. As a result, robustness and adaptability remain major challenges in the field.

To overcome these limitations, recent research has increasingly moved toward hybrid and explainable artificial intelligence approaches. Combining contextual text understanding with structural feature analysis, along with rule-based decision strategies, has shown promising results in improving both detection accuracy and model interpretability [9]. These approaches are especially useful in cybersecurity settings, where it is important not only to classify threats accurately but also to provide reasoning that can support analyst decision-making.

In this project, a hybrid approach is adopted that integrates transformer-based email analysis with URL intelligence and joint detection strategies. The proposed system builds on existing research while specifically addressing important limitations such as

cross-source robustness, feature portability, and practical applicability in real-world cybersecurity environments.

2.1 Traditional Machine Learning Approaches

Early phishing detection systems were largely based on classical machine learning algorithms such as Naïve Bayes, Decision Trees, and Support Vector Machines. These models became widely used in email classification tasks because they were relatively simple to implement, computationally efficient, and capable of producing useful results on structured datasets. Metsis et al. [1], for instance, demonstrated the effectiveness of Naïve Bayes in spam filtering, showing that probabilistic models could successfully identify unwanted emails on the basis of keyword frequency and distribution patterns.

Despite their usefulness, these traditional approaches rely heavily on manually engineered features and predefined decision rules. Their effectiveness therefore depends on the quality of the selected features and the extent to which those features reflect real phishing behavior. As phishing attacks became more sophisticated, attackers increasingly adopted obfuscation strategies designed to bypass keyword-based and rule-driven detection methods. Consequently, the performance of traditional models declined, particularly when applied to more advanced or carefully disguised phishing attempts.

2.2 Deep Learning-Based Approaches

With the rapid development of artificial intelligence, deep learning techniques have become increasingly important in phishing detection research. Unlike traditional machine learning models, deep learning methods are capable of learning complex patterns directly from data, reducing reliance on manual feature engineering. Transformer-based architectures such as BERT [6] marked a significant advancement by introducing contextual language understanding, allowing models to interpret meaning based on surrounding text rather than focusing only on isolated keywords.

DistilBERT [7], a compressed and computationally efficient version of BERT, has further improved the practicality of transformer-based methods by offering comparable

performance at lower computational cost. This makes it particularly suitable for applications that require faster inference or more efficient deployment. In addition to transformer models, architectures such as Convolutional Neural Networks (CNNs) and Long Short-Term Memory (LSTM) networks have also been used in phishing detection tasks [7]. These models are capable of identifying sequential and structural patterns in email content, enabling the detection of subtle indicators that may be overlooked by traditional classifiers.

Although deep learning approaches provide clear advantages in representation learning and contextual understanding, they also introduce practical challenges. In particular, they often require large volumes of labeled data and substantial computational resources for training and deployment. These requirements can make their use difficult in environments with limited infrastructure or restricted processing capacity.

2.3 URL-Based Phishing Detection

URL analysis represents another essential dimension of phishing detection, since malicious links frequently serve as the primary mechanism for redirecting victims to fraudulent websites. A large body of research has focused on identifying phishing URLs by examining their structural, lexical, and domain-related characteristics. Commonly used indicators include URL length, the presence of suspicious keywords, the use of direct IP addresses, unusual domain composition, and irregular formatting patterns.

Datasets such as PhishTank for malicious URLs and Top-1M websites for benign URLs are widely used in research for model training and evaluation. Studies that combine machine learning techniques with URL feature analysis [8] have reported encouraging results, demonstrating that structural and lexical properties can be effective in distinguishing malicious links from legitimate ones.

However, URL-based detection alone is not sufficient to address the full complexity of phishing attacks. Attackers increasingly use legitimate-looking domains, domain spoofing, shortened links, or even compromised trusted websites to evade detection mechanisms. As a result, URL analysis remains valuable, but it is most effective when

integrated with additional sources of contextual information.

2.4 Hybrid Phishing Detection Systems

To overcome the limitations of single-model approaches, researchers have increasingly proposed hybrid phishing detection systems that combine multiple detection strategies within a unified framework. Prior studies [3][4][8] have shown that integrating email content analysis with URL-based features can significantly improve detection accuracy and overall system robustness.

Hybrid systems are effective because they take advantage of complementary strengths. Email content analysis provides contextual understanding of language, tone, and intent, while URL analysis contributes structural and domain-level insights that may indicate malicious behavior. When these sources of information are combined, the system gains a more comprehensive view of the threat.

In addition to machine learning models, many hybrid approaches incorporate rule-based mechanisms such as domain validation, brand-domain mismatch detection, and heuristic scoring. These techniques further enhance system performance by introducing expert-driven checks that can capture suspicious behavior not always learned directly from data. As a result, hybrid systems are generally more robust and more suitable for real-world phishing detection, where attacks often involve multiple layers of deception.

2.5 Challenges and Research Gaps

Although considerable progress has been made in phishing detection, several important challenges remain unresolved.

One major issue is the lack of cross-domain generalization. Models that achieve strong performance on one dataset often perform poorly when evaluated on unseen datasets from different sources. This suggests that many systems learn source-specific patterns rather than broadly transferable phishing characteristics.

Another challenge is overfitting to specific patterns. A number of models depend heavily

on dataset-specific features, limiting their ability to adapt to new and evolving attack forms. This becomes particularly problematic in cybersecurity, where attacker behavior is dynamic rather than static.

The continuous evolution of phishing techniques presents an additional difficulty. Attackers regularly refine their methods in order to evade detection, using more convincing language, better visual imitation, and increasingly deceptive URL structures. Detection systems must therefore be adaptable enough to respond to emerging threats rather than only previously observed patterns.

A further concern is limited explainability. Many deep learning models operate as black-box systems, making it difficult for cybersecurity analysts to understand why a particular email or URL has been classified as suspicious. In practical environments, interpretability is important because analysts often require transparent reasoning to validate alerts and make informed security decisions.

Taken together, these challenges highlight the need for phishing detection systems that are not only accurate, but also robust, adaptable, and explainable.

2.6 Summary of Literature Review

The literature survey shows that both traditional machine learning and deep learning approaches have made important contributions to phishing detection, yet each category has its own limitations. Traditional models offer efficiency and simplicity but often struggle with complex contextual patterns, while deep learning models provide stronger representational power at the cost of higher computational requirements and reduced interpretability.

Recent studies suggest that hybrid approaches, which combine multiple detection techniques, offer superior performance and improved robustness compared to single-method systems. These approaches are particularly promising in phishing detection because they allow the system to incorporate both contextual understanding and structural analysis.

Based on the insights gathered from the literature, this project adopts a hybrid detection strategy that integrates transformer-based email analysis with URL intelligence and rule-based decision mechanisms. The goal of this approach is to improve cross-domain performance, enhance detection accuracy, and provide explainable results that are better suited for practical deployment in real-world cybersecurity environments.

Approach	Method	Strengths	Limitations	References
 Traditional Machine Learning	Uses algorithms like Naive Bayes, SVM, Decision Trees with handcrafted features.	• Fast training • Simple to implement • Works well on structured data	• Manual feature engineering • Limited for complex patterns • Lower generalization	[10], [11], [12]
 Deep Learning Approaches	Uses models like CNN, RNN, LSTM, Transformers to learn features automatically.	• High accuracy • Learns complex patterns • Effective for large datasets	• Requires large datasets • High computational resources • Less interpretable	[3], [4], [7]
 Hybrid Approaches	Combines ML + DL, content analysis, URL analysis, and rule-based methods.	• Better detection accuracy • Balances precision and recall • Handles diverse attacks	• More complex implementation • Higher development time • Requires careful integration	[2], [8], [9]
 Rule-Based Methods	Relies on predefined rules, blacklists, heuristics, and signatures.	• Easy to implement • Works without training data • Effective against known threats	• Cannot detect new or obfuscated attacks • Requires frequent rule updates • High false positives	[4], [6]

Figure 2.1: Comparison of phishing detection approaches

CHAPTER – 3

PROBLEM STATEMENT

Phishing attacks have become increasingly sophisticated, often combining deceptive email content with malicious URLs in order to manipulate users and bypass conventional security mechanisms. Traditional phishing detection systems typically rely on a single mode of analysis, such as email content filtering or URL-based detection. Although these approaches may perform reasonably well under controlled conditions, they are often less effective in real-world environments where attackers employ multiple, layered techniques to evade detection.

One of the major limitations of existing phishing detection systems is their weak ability to generalize across different datasets and evolving attack patterns. Models trained on a specific dataset frequently overfit to source-specific characteristics and show a noticeable decline in performance when tested on previously unseen data. This limitation is particularly serious in phishing detection because attacker strategies are not static; they continuously evolve in response to emerging defenses, reducing the long-term effectiveness of rigid or narrowly trained models.

Another important challenge lies in maintaining an effective balance between precision and recall. Systems designed to maximize detection rates often generate a large number of false positives, which can overwhelm users as well as cybersecurity analysts and reduce trust in the system. Conversely, systems that prioritize reducing false positives may become too conservative and fail to identify subtle phishing attempts, thereby increasing the risk of successful attacks and security breaches.

In addition, many advanced machine learning models suffer from limited interpretability, making it difficult for users and analysts to understand the reasoning behind a particular prediction. This reduces their practical value in real-world cybersecurity environments, where explainability is essential for validating alerts, supporting decision-making, and improving user confidence in automated systems.

Therefore, there is a clear need for a phishing detection system that is robust, scalable, and capable of addressing the limitations of existing approaches. Such a system should be able to integrate multiple sources of information, generalize across diverse datasets, maintain a practical balance between precision and recall, and provide outputs that are sufficiently explainable for operational use. The objective of this project is to address these challenges by developing a hybrid phishing detection system that combines email content analysis, URL intelligence, and rule-based decision mechanisms within a unified framework.

CHAPTER – 4

RATIONALE / NEED FOR THE PROJECT

The rapid expansion of digital communication, online services, and internet-based transactions has significantly increased exposure to cyber threats, among which phishing has emerged as one of the most prevalent and damaging attack methods. Phishing attacks target both individuals and organizations by exploiting human trust through deceptive emails, fraudulent messages, and malicious links, often resulting in financial loss, data breaches, identity theft, and unauthorized access to sensitive systems.

Despite the availability of various phishing detection techniques, many existing systems remain limited in their effectiveness because they rely on a single mode of analysis. Systems that focus only on email content may fail when attackers use carefully crafted language that appears legitimate and avoids obvious indicators of fraud. Similarly, URL-based systems may overlook threats when malicious intent is hidden within links that closely resemble trusted domains or exploit compromised legitimate websites. This creates an important gap in detection capability, particularly in real-world environments where attackers combine multiple deceptive techniques to bypass security controls.

Another major concern is the limited adaptability of traditional machine learning models. Many existing systems are trained on narrow or source-specific datasets and do not generalize well across different data distributions. As a result, their performance often declines when they are exposed to unseen datasets or emerging phishing strategies. In a threat landscape where attacker behavior evolves continuously, such limitations reduce the reliability and long-term usefulness of static detection models.

In addition to detection performance, interpretability has become an increasingly important requirement in practical cybersecurity systems. Security analysts and end users need to understand why a particular email or URL has been classified as suspicious in order to make informed decisions and respond appropriately. When models operate as black boxes and provide little insight into their predictions, user trust is reduced and adoption in real-world settings becomes more difficult.

The rationale behind this project is therefore to develop a more comprehensive, reliable, and practically useful phishing detection system that addresses these limitations. By integrating email content analysis, URL intelligence, and rule-based decision mechanisms within a unified framework, the proposed system seeks to improve detection accuracy, strengthen cross-domain robustness, and provide more meaningful insight into prediction outcomes.

This project is particularly relevant in modern cybersecurity environments, where organizations require scalable, efficient, and adaptive solutions to defend against increasingly sophisticated attacks. In addition to improving technical detection capability, the proposed system places emphasis on usability, interpretability, and practical deployment, making it a valuable and timely contribution to the field of phishing detection.

4.1 Increasing Threat of Phishing Attacks

Phishing attacks have increased significantly alongside the expansion of digital platforms, online banking, e-commerce, and cloud-based services. As more personal, financial, and organizational activities move into digital environments, attackers have gained greater opportunities to exploit users through deceptive communication. Modern phishing campaigns are no longer limited to simple fraudulent emails; they increasingly employ advanced techniques such as social engineering, domain spoofing, impersonation, and carefully crafted deceptive content to mislead victims.

These attacks affect not only individual users but also organizations of all sizes, often resulting in financial loss, data compromise, reputational damage, and unauthorized access to critical systems. The growing scale and sophistication of these threats underline the need for phishing detection systems that are more accurate, adaptive, and reliable than conventional solutions.

4.2 Limitations of Existing Systems

Most existing phishing detection systems rely on either email content analysis or URL-based detection as separate and independent approaches. While these methods may perform well in controlled or narrowly defined environments, they often prove insufficient in real-world scenarios where attackers deliberately combine multiple deceptive techniques to evade detection.

A further limitation is that many existing models suffer from overfitting to the datasets on which they are trained. As a result, they may achieve strong performance during testing on known data but fail to maintain the same effectiveness when applied to unseen datasets or new phishing patterns. This lack of generalization significantly reduces the reliability and practical usefulness of traditional detection systems in dynamic cybersecurity environments.

4.3 Need for Hybrid Detection Approach

To address the limitations of single-mode detection systems, there is a clear need for a hybrid approach that integrates multiple sources of information within a unified framework. Combining email content analysis with URL intelligence enables the system to capture both contextual and structural indicators of phishing attacks.

Email analysis helps identify suspicious language, semantic intent, and deceptive communication patterns, while URL analysis provides insight into link structure, domain characteristics, and other indicators of malicious behavior. In addition, the inclusion of rule-based mechanisms can further strengthen detection performance by incorporating domain knowledge, heuristic reasoning, and expert-defined security checks. Together, these components support a more comprehensive and robust approach to phishing detection.

4.4 Importance of Explainable Systems

Modern cybersecurity solutions must provide not only accurate predictions but also

results that are understandable and interpretable to users. Security analysts and end users often need to know why a particular email or URL has been flagged as suspicious in order to make informed decisions and take appropriate action.

Models that operate as black boxes without offering any explanation can reduce user trust and limit practical adoption, especially in operational security settings where transparency is important. For this reason, explainability has become a critical requirement in phishing detection systems. Incorporating explainable outputs helps improve confidence in the system, supports analyst workflows, and enhances the overall usability of the solution in real-world deployment environments.

4.5 Practical Relevance and Applications

The proposed phishing detection system has strong practical relevance across a variety of real-world cybersecurity settings. It can be integrated into email filtering platforms, browser extensions, enterprise security gateways, and other defensive systems designed to identify and block phishing threats before they reach users.

Its ability to support multiple detection modes and provide explainable outputs makes it suitable for both individual users and organizational environments. For individual users, the system can serve as an additional layer of protection against fraudulent emails and malicious links. For organizations, it can support security operations by improving alert quality, reducing missed threats, and enhancing decision-making. These characteristics increase the overall usability, scalability, and deployment potential of the proposed system within modern cybersecurity frameworks.

CHAPTER – 5

OBJECTIVES

The objective of this project is to design and develop a comprehensive phishing detection system capable of identifying malicious emails and URLs with a high degree of accuracy, reliability, and practical usefulness. As phishing attacks continue to grow in sophistication, it has become necessary to move beyond traditional detection methods and adopt more advanced approaches that combine multiple sources of information within a unified framework.

This project focuses on the development of a hybrid detection framework that integrates contextual analysis of email content with structural analysis of URLs. By leveraging machine learning models, particularly transformer-based architectures such as DistilBERT, together with engineered feature extraction and rule-based decision mechanisms, the proposed system aims to provide a more robust, scalable, and adaptable solution for phishing detection.

In addition to improving detection performance, the project places strong emphasis on real-world applicability, cross-domain adaptability, and usability for both end users and cybersecurity analysts. For this reason, the objectives of the project are organized into primary, technical, system-level, and user-oriented categories in order to ensure a balanced and practically relevant implementation.

5.1 Primary Objectives

The primary objectives define the broad goals of the project and focus on the core purpose of the phishing detection system.

- To design and develop an efficient phishing detection system capable of accurately identifying malicious emails and URLs.
- To improve detection performance in real-world environments where phishing attacks are dynamic, diverse, and continuously evolving.
- To strengthen the generalization capability of the system across multiple datasets

and varied attack patterns.

- To reduce both false positives and false negatives in order to achieve a balanced, dependable, and effective detection mechanism.

5.2 Technical Objectives

The technical objectives focus on the implementation of the system and the use of advanced machine learning and feature engineering techniques.

- To implement a transformer-based model, specifically DistilBERT, for analyzing email subject lines and body content.
- To extract meaningful URL features, including structural, lexical, and behavioral characteristics.
- To incorporate engineered indicators such as URL length, suspicious lexical patterns, and domain inconsistencies.
- To combine contextual understanding from textual data with feature-based analysis in order to improve overall detection accuracy.
- To optimize model performance through suitable training strategies, evaluation metrics, and threshold tuning techniques.

5.3 System Objectives

The system objectives define how the major components of the project are integrated into a unified and operational framework.

- To develop a hybrid detection system that combines email analysis and URL intelligence within a single processing pipeline.
- To implement a joint detection mechanism that merges the outputs of multiple models to support more reliable classification.
- To incorporate rule-based decision logic, including domain validation and brand-domain mismatch detection.
- To design the system to support multiple operational modes, such as high-recall detection and high-confidence alerting.

- To ensure scalability and flexibility so that the system can be deployed in different practical environments.

5.4 User-Oriented Objectives

The user-oriented objectives focus on usability, interpretability, and practical value for real-world users and analysts.

- To provide clear and explainable outputs that help users understand the reasoning behind system predictions.
- To support batch analysis for efficient processing of multiple emails and URLs.
- To design an intuitive interface that improves usability, accessibility, and overall user experience.
- To ensure that the system can be effectively used by both individual users and cybersecurity analysts.
- To strengthen trust in automated phishing detection through transparency, consistency, and reliability.

5.5 Overall Objective Summary

The overall objective of this project is to develop a phishing detection system that is scalable, accurate, and practically applicable, while addressing the key limitations of traditional approaches. By combining advanced machine learning techniques with rule-based intelligence and a user-centered design perspective, the proposed system aims to provide a comprehensive and effective solution to the challenges of modern phishing detection.

CHAPTER – 6

FEASIBILITY STUDY

A feasibility study is an essential stage in system design because it helps determine whether a proposed solution can be developed and implemented successfully within the available resources, constraints, and technical conditions. It provides a structured assessment of the practicality and viability of a project before full-scale implementation begins.

In the context of phishing detection, such an evaluation is particularly important because the proposed system must function effectively in real-world environments while maintaining accuracy, scalability, efficiency, and usability. The feasibility of the proposed hybrid phishing detection system is therefore examined from multiple perspectives, including technical, economic, operational, implementation, and long-term sustainability considerations.

The proposed system is designed using modern machine learning techniques and open-source technologies in order to deliver a solution that is efficient, cost-effective, and scalable. The following subsections present a detailed discussion of the feasibility of the project.

6.1 Technical Feasibility

Technical feasibility refers to whether the required technologies, tools, datasets, and expertise are available to design, develop, and implement the proposed system effectively.

The phishing detection system is built using widely adopted and well-supported technologies such as Python for programming, FastAPI for backend development, and established machine learning libraries including Scikit-learn and HuggingFace Transformers. These technologies provide a strong development foundation and are suitable for building an intelligent detection framework. The use of DistilBERT further supports efficient contextual analysis of email content while maintaining a practical balance between model performance and computational cost.

The datasets used for training and evaluation, including the Enron dataset, SpamAssassin corpus, Nazario phishing dataset, PhishTank URLs, and Top-1M benign URLs, are publicly available and commonly used in research. Their accessibility supports both reproducibility and ease of experimentation.

In addition, the system is designed to operate efficiently on standard computing environments, including CPU-based systems. Although GPU acceleration may improve the speed of model training, it is not essential for deployment. This makes the system practical even in resource-constrained settings.

The modular architecture of the proposed framework also strengthens its technical feasibility by allowing new models, features, or detection rules to be incorporated with minimal disruption. For these reasons, the project is considered technically feasible and well aligned with current technological capabilities.

6.2 Economic Feasibility

Economic feasibility examines the cost implications associated with the development, testing, and deployment of the proposed system.

The project is highly cost-effective because it depends primarily on open-source software and publicly available datasets. Programming languages such as Python, along with frameworks such as FastAPI, Scikit-learn, and HuggingFace Transformers, can be used without licensing costs. This substantially reduces the financial burden associated with system development.

The hardware requirements are also relatively modest. The system can be developed, tested, and deployed on standard personal computers or commonly available computing environments, which removes the need for expensive infrastructure or specialized hardware in many cases.

Beyond development costs, the system also offers meaningful long-term economic benefits. Effective phishing detection can help prevent financial fraud, credential

compromise, data breaches, and operational disruption. The cost savings associated with reducing such incidents can significantly outweigh the comparatively low cost of implementation.

Therefore, from an economic perspective, the project is feasible and offers strong value in relation to its cybersecurity benefits.

6.3 Operational Feasibility

Operational feasibility evaluates how well the system can function in real-world settings and how easily it can be adopted and used by its intended users.

The proposed phishing detection system is designed with practical deployment and usability in mind. It supports multiple operational modes, such as high-recall detection for continuous monitoring environments and high-confidence detection for more sensitive or critical use cases. This flexibility allows the system to be adapted to different security requirements and application contexts.

An important operational strength of the system is the inclusion of explainable outputs. Rather than generating predictions without justification, the system is intended to provide insight into why a particular email or URL has been classified as phishing or benign. This is especially valuable for cybersecurity analysts, who must interpret system outputs and make informed security decisions.

The system is also suitable for integration into multiple deployment environments, including email filtering platforms, browser extensions, and enterprise security infrastructures. Its modular design further improves operational feasibility by enabling customization and extension according to organizational needs.

Considering its flexibility, usability, interpretability, and integration potential, the proposed system is operationally feasible for practical cybersecurity applications.

6.4 Implementation Feasibility

Implementation feasibility considers whether the project can realistically be completed within the available time, resources, and skill constraints.

The project follows a structured and manageable development process consisting of data collection, preprocessing, feature extraction, model training, evaluation, and deployment. Each of these stages can be carried out using available tools, established workflows, and existing technical knowledge.

The implementation process benefits from the use of well-documented libraries and standard machine learning practices, which reduces development complexity and improves efficiency. In addition, the modular design of the system allows different components to be developed, tested, and refined independently, making the workflow more manageable.

The expected timeline for the project is realistic within an academic setting, and the required technical skills—such as programming, machine learning, and basic web development—are attainable within the project scope. Consequently, the implementation of the proposed system is feasible within the available academic and technical constraints.

6.5 Scalability and Maintainability Feasibility

Scalability and maintainability are critical factors in assessing the long-term practicality of any cybersecurity system.

The proposed phishing detection system is designed with scalability in mind, allowing it to handle increasing volumes of data and user requests without substantial loss of performance. The use of relatively efficient models such as DistilBERT supports this objective by providing strong contextual analysis without excessive computational overhead.

Maintainability is supported through the modular architecture of the framework. Individual components, such as the email analysis model, URL analysis module, and

rule-based decision engine, can be updated, replaced, or improved independently. This design simplifies future development and reduces the difficulty of adapting the system to new requirements.

Such flexibility is especially important in phishing detection, where attack techniques evolve continuously. A maintainable and scalable system is better positioned to respond to emerging threats, incorporate new datasets, and integrate improved detection strategies over time. For these reasons, the proposed system demonstrates strong long-term viability.

6.6 Summary of Feasibility

Based on the assessment of technical, economic, operational, implementation, and scalability considerations, the proposed phishing detection system can be regarded as highly feasible. The required technologies are readily available, the cost of development is relatively low, the implementation process is manageable, and the system is well suited for practical deployment in real-world environments.

In addition to meeting current project requirements, the proposed framework also provides a strong foundation for future improvements and extensions. This ensures that the system remains relevant and adaptable within an evolving cybersecurity landscape, thereby reinforcing its value as a practical and effective phishing detection solution.

Aspect	Feasibility Status	Remarks	Icon
 Technical	High	The project uses modern machine learning tools and frameworks that are widely supported and technically achievable.	
 Economic	High	The system is cost-effective as it relies on open-source tools and public datasets, minimizing development and deployment costs.	
 Operational	High	The system is user-friendly and can be easily integrated into different platforms and real-world environments.	
 Implementation	High	The project is achievable within the given timeline and resource constraints using available skills and tools.	

Figure 6.1: Feasibility summary

CHAPTER – 7

METHODOLOGY

The methodology of the proposed phishing detection system is centered on the development of a hybrid framework that integrates email content analysis with URL-based intelligence. The system is designed to address the limitations of traditional single-model approaches by combining multiple detection strategies within a unified processing pipeline.

The overall methodology consists of several key stages, including data preprocessing, feature extraction, model training, and a joint detection mechanism that produces the final classification decision. Each stage contributes to improving the accuracy, robustness, and practical effectiveness of the system.

7.1 System Overview

The proposed system follows a multi-stage pipeline in which input data is processed through multiple analytical components before a final phishing classification is generated. The input to the system includes email content, specifically the subject and body, together with any URLs embedded within the email.

The system is organized into three primary components:

- Email analysis module
- URL analysis module
- Joint detection mechanism

These components work together to evaluate both the contextual and structural characteristics commonly associated with phishing attacks. By combining these complementary perspectives, the system is able to provide a more comprehensive assessment than methods that rely on a single source of information.

7.2 Data Collection and Preprocessing

The proposed system is trained and evaluated using multiple publicly available datasets in order to ensure diversity of data and improve generalization across different sources. The datasets used in this study include:

- Enron dataset
- SpamAssassin dataset
- Nazario phishing corpus
- PhishTank URL dataset
- Top-1M benign URL dataset

The collected data is preprocessed to improve consistency, quality, and suitability for model training. Email text is cleaned by removing unnecessary characters, normalizing formatting, and combining the subject and body into a unified textual representation. URLs are extracted from email content and processed separately for structural feature analysis.

This preprocessing stage is essential because it reduces noise, standardizes the data format, and ensures that both textual and URL inputs are prepared appropriately for subsequent machine learning and feature extraction procedures.

7.3 Email Content Analysis

The email analysis module is responsible for identifying phishing-related patterns within the textual content of an email. For this purpose, a transformer-based model, DistilBERT, is employed in order to capture contextual and semantic relationships present in the text.

The input to the email analysis model consists of:

- Email subject
- Email body

These elements are combined into a single input sequence and then tokenized before

being passed to the model. DistilBERT generates contextual embeddings that represent the meaning and structure of the text, and these embeddings are subsequently used for classification.

This approach enables the system to detect phishing attempts even when attackers avoid explicit keywords and instead rely on subtle wording, persuasive language, or context-dependent deception. As a result, the model is better equipped to identify sophisticated phishing emails than traditional keyword-based methods.

7.4 URL Feature Analysis

The URL analysis module is designed to identify malicious characteristics in the links embedded within emails. Since phishing campaigns often rely on deceptive URLs to redirect victims to fraudulent websites, URL analysis serves as an important complement to email text analysis.

A range of structural and lexical features is extracted from each URL, including:

- URL length
- Presence of suspicious keywords
- Use of IP addresses instead of domain names
- Domain mismatch indicators
- Number of URLs contained within the email

These features are used to estimate the likelihood that a URL is malicious. Unlike the email analysis module, which focuses on contextual language patterns, the URL analysis component concentrates on structural properties and link-related indicators. Together, these two perspectives strengthen the overall detection capability of the system.

7.5 Feature Engineering

Feature engineering plays an important role in improving the performance of the proposed phishing detection system. In addition to raw textual and URL data, several supplementary features are derived from email content and embedded links in order to

capture patterns that may not be fully represented through raw input alone.

Examples of engineered features include:

- Digit ratio and capitalization ratio in email text
- Number of exclamation marks
- Presence of shortened URLs
- Suspicious top-level domains

These engineered features help the system capture stylistic, structural, and heuristic indicators that are often associated with phishing behavior. Their inclusion enhances the model's ability to identify malicious patterns that may otherwise remain undetected.

7.6 Joint Detection Mechanism

The joint detection mechanism combines the outputs of the email analysis and URL analysis modules in order to produce a final phishing classification. This stage is central to the hybrid nature of the proposed system, as it integrates evidence from multiple sources before making a final decision.

The combination process is performed using:

- Weighted scoring
- Rule-based optimization
- Threshold-based decision-making

Important rule-based strategies include:

- Brand-domain mismatch detection
- Suspicious URL escalation
- Trusted domain validation

By incorporating these mechanisms, the system is able to use both learned model predictions and expert-defined decision rules. This hybrid strategy improves overall robustness and allows the system to respond more effectively to complex phishing attempts that may not be captured by a single analytical component.



Figure 7.1: Workflow architecture of Joint detection Engine.

7.7 Threshold Optimization

Threshold tuning is applied to adjust the balance between precision and recall according to different operational requirements. Since phishing detection systems may be used in varied deployment settings, a single fixed threshold is not always appropriate.

The proposed system therefore supports multiple operational modes:

- **High-recall mode:** prioritizes the detection of as many phishing attempts as

- possible and is suitable for security monitoring environments.
- **Balanced mode:** maintains a practical trade-off between precision and recall for general-purpose deployment.
 - **High-confidence mode:** reduces false positives and is more suitable for critical environments where only highly reliable alerts are desired.

This flexibility allows the system to be tailored to different application scenarios while maintaining effective detection performance.

7.8 System Workflow

The overall workflow of the proposed system can be summarized as follows:

- Input email content and embedded URLs are collected.
- Data preprocessing is performed to clean and standardize the inputs.
- Email content is analyzed using the DistilBERT-based model.
- URL features are extracted and evaluated through the URL analysis module.
- Outputs from both modules are combined.
- Rule-based optimization is applied to refine the decision.
- A final phishing classification is generated.

This workflow reflects the integration of contextual analysis, structural evaluation, and decision-level optimization within a single coherent detection framework.

7.9 Summary of Methodology

The proposed methodology combines advanced machine learning techniques with rule-based decision mechanisms to develop a robust and practical phishing detection system. By integrating contextual understanding from email content with structural analysis of URLs, the system is able to achieve improved detection accuracy, stronger robustness, and better generalization across different datasets.

Overall, the methodology provides a comprehensive foundation for building a phishing detection framework that is not only technically effective but also adaptable to real-world cybersecurity requirements.




Figure 7.2: Workflow diagram of final system.

Chapter – 8

Project Implementation Plan/Timeline

The development of the proposed phishing detection system was carried out in a structured and systematic manner over a defined project period. To ensure smooth progress and effective management of tasks, the work was divided into multiple phases, with each phase focusing on a specific aspect of system development. These phases included background research, data preparation, model development, system integration, performance evaluation, and deployment.

Establishing a clear timeline was important for maintaining direction throughout the project and ensuring that each stage was completed within the available time frame. This phased approach also helped in organizing the workflow efficiently, identifying dependencies between tasks, and allowing sufficient time for testing and refinement. The following section outlines the major development phases of the project.

8.1 Development Phases

The project was executed through the following development phases:

Phase 1: Research and Problem Understanding

This phase focused on studying the phishing detection domain, understanding the major challenges associated with modern phishing attacks, and reviewing existing methods proposed in the literature. It provided the conceptual foundation for defining the problem and selecting an appropriate system design.

Phase 2: Data Collection and Preprocessing

During this phase, relevant datasets were collected from publicly available sources, and the data was prepared for experimentation. This included cleaning the email content, extracting URLs, normalizing inputs, and organizing the datasets for model training and evaluation.

Phase 3: Model Development

This phase involved the implementation of the phishing detection models. Email analysis and URL detection components were developed using suitable machine learning techniques, including transformer-based methods for textual analysis and feature-based methods for URL evaluation.

Phase 4: Hybrid System Integration

In this phase, the individual detection modules were combined into a unified hybrid framework. A joint detection mechanism was developed to integrate multiple signals and support more reliable phishing classification.

Phase 5: Evaluation and Optimization

This phase focused on assessing system performance using appropriate evaluation metrics. Threshold tuning, error analysis, and performance optimization were carried out to improve accuracy, robustness, and overall reliability.

Phase 6: Deployment and Testing

The final phase involved integrating the developed models into an application-level environment and testing the system under realistic usage conditions. This stage helped evaluate the practical applicability of the solution and its readiness for real-world deployment scenarios.

8.2 Project Timeline Table

The following table presents the month-wise distribution of major project activities carried out during the development of the phishing detection system.

Month	Activity
January	Literature Survey and Problem Statement
February	Data Collection, Preprocessing and Model Development
March	Integration of Hybrid Detection System and Optimization
April	Final Report Preparation and Submission

Table 8.1: Project Timeline Table

8.3 Summary of Timeline

The structured timeline provided a clear and systematic framework for carrying out the project in an organized manner. It supported efficient task management, ensured that each development phase was completed within the planned duration, and helped maintain steady progress throughout the project.

This phased approach also allowed room for iterative refinement, enabling improvements to be made at different stages on the basis of experimental findings and evaluation results. As a result, the timeline contributed not only to the timely completion of the project, but also to the overall quality and consistency of the final system.

CHAPTER – 9

RESULTS AND DISCUSSION

This chapter presents the evaluation results of the proposed phishing detection system and provides an analysis of its performance across different models and configurations. The findings demonstrate the effectiveness of the hybrid framework in improving both detection accuracy and robustness when compared with individual detection components.

The system was evaluated using multiple datasets containing both phishing and legitimate emails and URLs. Performance was assessed using standard classification metrics, including accuracy, precision, recall, F1-score, and ROC-AUC. These metrics provide a broad understanding of how well the system performs under different operating conditions.

9.1 Evaluation Metrics

To evaluate the performance of the proposed system, the following standard metrics were used:

- **Accuracy:** Measures the overall proportion of correct predictions made by the system.
- **Precision:** Measures the proportion of instances predicted as phishing that are actually phishing.
- **Recall:** Measures the proportion of actual phishing instances that are correctly detected by the system.
- **F1-score:** Represents the harmonic mean of precision and recall and provides a balanced measure of classification performance.
- **ROC-AUC:** Measures the ability of the model to distinguish between phishing and benign instances across different threshold settings.

Together, these metrics provide a comprehensive and reliable basis for evaluating the effectiveness of the phishing detection system.

9.2 Email Model Performance

The email-based detection model was first evaluated independently using the test datasets. The results are presented below.

Metric	Value
Accuracy	0.4643
Precision	0.4706
Recall	0.8889
F1-Score	0.6154
ROC-AUC	0.5939

Table 9.1: Email Model Performance

Analysis

The email model achieved a high recall, indicating that it was effective in detecting a large proportion of phishing emails. This suggests that the transformer-based text analysis component was able to identify phishing-related contextual patterns even when the language was not obviously suspicious. However, the relatively low precision indicates that the model also produced a considerable number of false positives. In other words, although the model was sensitive to phishing indicators, it was less reliable when used in isolation and showed limitations in terms of generalization and classification balance.

9.3 URL Model Performance

The URL-based detection model was also evaluated independently in order to examine its contribution to the overall system.

Metric	Value
Accuracy	0.6607

Precision	0.8333
Recall	0.3704
F1-Score	0.5128
ROC-AUC	0.7752

Table 9.2: URL Model Performance

Analysis

The URL model produced high precision, which indicates that when it classified a link as phishing, that prediction was usually correct. This shows that structural and lexical URL features were effective in confirming malicious behavior. However, the recall was relatively low, meaning that a substantial number of phishing URLs were missed. This suggests that the URL model was stronger as a confirmation mechanism than as a discovery mechanism. While it was effective at reducing false alarms, it was less capable of identifying all malicious cases on its own.

9.4 Joint Model Performance

The hybrid detection system was then evaluated by combining both email and URL analysis within a unified decision framework.

Metric	Value
Accuracy	0.7321
Precision	0.7188
Recall	0.7931

F1-Score	0.7541
ROC-AUC	0.7241

Table 9.3: Joint Model Performance

Analysis

The joint model achieved a more balanced performance across all evaluation metrics when compared with the individual models. In particular, the F1-score improved noticeably, reflecting a stronger balance between precision and recall. These results demonstrate the practical value of the hybrid approach, as the combined model was able to benefit from the strengths of both the email and URL components. The email model contributed stronger recall, while the URL model improved precision, resulting in a more reliable overall system.

Main Detection Workspace
Configure analysis mode, threshold profile, and strategy

URL Detection
Analyze standalone URLs

Email Detection
Analyze message intent

Joint Detection
Fuse email + URL signals

OPERATING MODE
Balanced - Default

JOINT STRATEGY
Optimized (rule-assisted)

Email Subject
Urgent: Verify Your Account to Avoid Suspension

Email Body
We noticed unusual activity on your account and have temporarily restricted access for your protection. To restore full functionality, please verify your identity immediately.

Failure to complete verification within 24 hours may result in permanent suspension.

http://secure-account-verify-login.com/update
http://192.168.45.12/bank/verify

Sender (optional)
alerts@company.com

Manual URLs (optional)
Add comma/newline-separated URLs if not present in body

Figure 9.1: Workspace Input Window

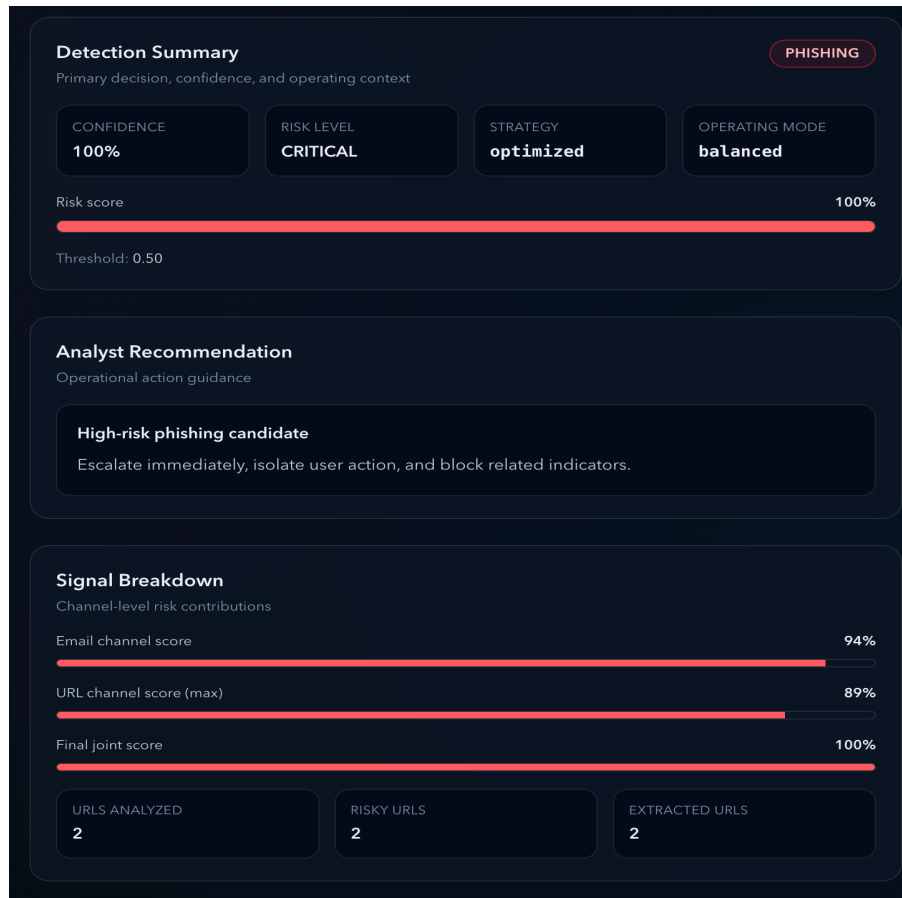


Figure 9.2: Workspace Output Window

9.5 Optimized Joint Model Performance (Final System)

After further optimization and expansion of the evaluation dataset, the final version of the joint detection system achieved the following results:

Metric	Value	F1-Score
Balanced(0.5)	0.925	0.939
High Confidence(0.6)	0.945	0.951

Table 9.4: Optimized Joint Model Performance

Analysis

The optimized joint system showed a substantial improvement in performance after refinement. Achieving accuracy above 90% in the expanded evaluation setting indicates that the hybrid framework became significantly more stable and effective after optimization. The high-confidence mode further improved reliability by reducing false positives, which is especially important in critical deployment environments where excessive alerts can reduce trust in the system. These results suggest that the optimized system has strong potential for practical real-world use.




Figure 9.3: Accuracy comparison graph

9.6 Confusion Matrix Analysis

The confusion matrix provides a more detailed view of classification outcomes by showing how predictions are distributed across the four possible categories:

- **True Positives:** phishing instances correctly identified as phishing
- **True Negatives:** benign instances correctly identified as benign

- **False Positives:** benign instances incorrectly classified as phishing
- **False Negatives:** phishing instances incorrectly classified as benign

Insight

The confusion matrix analysis indicates that the initial individual models suffered from imbalance-related issues. The email model tended to produce more false positives, while the URL model missed a larger number of phishing cases. By contrast, the hybrid system reduced false negatives through improved recall and, after optimization, also reduced false positives through better decision thresholds and rule refinement. This shows that the integration of multiple detection strategies helped create a more stable and balanced classification outcome.

9.7 Error Analysis

Further analysis was conducted to better understand the types of errors made by the system and the situations in which misclassification occurred.

Common False Positives:

- Legitimate banking alerts
- Travel confirmation emails
- Emails containing multiple links

Common False Negatives:

- Subtle phishing emails with persuasive but non-obvious language
- Malicious URLs designed to appear clean and trustworthy

Insight

These errors highlight the complexity of phishing detection and reinforce the importance of combining multiple analytical perspectives. In particular, they show the need for stronger contextual understanding, more reliable domain validation, and continued refinement of hybrid detection logic. The findings also suggest that even effective models require ongoing improvement to address edge cases and increasingly deceptive phishing strategies.

9.8 Discussion

The experimental results clearly demonstrate that single-model approaches are insufficient for robust phishing detection in realistic settings. The email model provided strong recall and was effective in identifying many phishing emails, but it also generated a relatively high number of false positives. The URL model, on the other hand, offered strong precision and was more reliable when confirming phishing cases, yet it missed many malicious instances when used independently.

The hybrid model successfully balanced these strengths and weaknesses by combining contextual text understanding with structural URL analysis. As a result, it achieved stronger overall performance than either individual component. The optimized joint detection system further improved both accuracy and generalization, indicating that hybrid decision-making is a more effective strategy for handling the complexity of modern phishing attacks.

These findings support the broader argument of this project: phishing detection systems should not rely on a single source of information, but instead should integrate complementary signals to achieve better robustness, adaptability, and practical deployment value.

9.9 Summary of Results

The results of this study show that the proposed phishing detection system achieves high accuracy and strong robustness by combining multiple detection strategies within a hybrid framework. The joint model outperformed the individual email and URL models by achieving a better balance between precision and recall, while the optimized final system demonstrated strong performance suitable for real-world use.

Overall, the findings confirm that the integration of contextual email analysis, URL intelligence, and decision-level optimization provides a practical and effective solution for modern phishing detection challenges.

CHAPTER – 10

CONCLUSION

The growing sophistication of phishing attacks has made the development of advanced detection systems increasingly necessary. As attackers continue to refine their methods, phishing campaigns are no longer limited to easily identifiable fraud but often involve subtle language, deceptive domain structures, and coordinated evasion strategies. In such conditions, traditional approaches that rely only on email content analysis or URL-based detection are often insufficient for effective real-world protection.

This project presented the design and development of a hybrid phishing detection system that integrates email content analysis with URL intelligence within a unified framework. The proposed system employs transformer-based models, particularly DistilBERT, to capture contextual and semantic information from email text, while also incorporating feature-based and rule-assisted analysis of URLs to detect structural indicators of malicious intent. By combining these complementary sources of information, the system is able to identify both linguistic and structural characteristics commonly associated with phishing attacks.

The performance evaluation shows that while the individual detection models each possess distinct strengths and limitations, the hybrid detection mechanism offers a substantial improvement in overall accuracy and robustness. In particular, the optimized joint model achieved strong performance, exceeding 90% accuracy on an expanded dataset, while also maintaining a more balanced trade-off between precision and recall. These findings demonstrate that combining multiple analytical perspectives leads to more reliable phishing detection than relying on isolated models.

An additional strength of the proposed system lies in its flexibility and practical usability. The system supports multiple operational modes, allowing it to be adapted to different deployment requirements, such as high-recall monitoring for security operations and high-confidence alert generation for more critical environments. Furthermore, the

inclusion of explainable outputs improves transparency and user trust, making the system more useful not only for automated filtering but also for cybersecurity analysts who require interpretable decision support.

Overall, the proposed system successfully addresses several important limitations of traditional phishing detection methods and provides a scalable, efficient, and practically relevant solution. The project demonstrates that the integration of machine learning models with rule-based intelligence can significantly strengthen phishing detection capabilities and contribute meaningfully to modern cybersecurity defense in increasingly complex digital environments.

CHAPTER – 11

FUTURE SCOPE

Although the proposed phishing detection system demonstrates strong performance and achieves high accuracy, there remain several opportunities for further improvement and expansion. Because cyber threats continue to evolve in complexity and sophistication, phishing detection systems must also advance in order to remain effective against emerging attack techniques. Future development of this project can therefore focus on enhancing model capability, expanding deployment potential, and improving adaptability in real-world cybersecurity environments.

The following areas represent important directions for future work.

11.1 Advanced Model Improvements

One important direction for future enhancement is the use of more advanced language models and learning strategies. Larger transformer architectures such as BERT or RoBERTa may provide stronger contextual understanding and richer semantic representations, which could improve detection accuracy for subtle and highly deceptive phishing messages.

In addition, ensemble learning techniques could be explored to combine the strengths of multiple machine learning models. Such an approach may improve classification stability and reduce dependence on the limitations of any single model. Another promising direction is continuous model updating through training on real-time or newly collected data, which would allow the system to adapt more effectively to evolving phishing strategies.

11.2 Real-Time Threat Detection

A major area for future development is the extension of the system toward real-time phishing detection. This would enable the system to analyze emails and URLs dynamically as they are received, making it more suitable for operational cybersecurity settings.

Future implementations could include integration into live email filtering systems, browser extensions capable of detecting malicious links instantly, and enterprise-level security platforms for continuous monitoring. Real-time capability would significantly increase the practical value of the system in active defense environments.

11.3 Improved Feature Engineering

Further improvements can also be achieved through more advanced feature engineering. In addition to the features already used in the current system, future versions could incorporate domain reputation scores, WHOIS-based information, and other external threat intelligence signals to strengthen detection reliability.

The system could also be extended to detect homograph attacks, domain impersonation, and more sophisticated forms of URL obfuscation. Behavioral features may also be incorporated to identify suspicious interaction patterns, such as unusual click behavior or abnormal navigation sequences. These additions could help the system capture more subtle and evolving phishing indicators.

11.4 System Integration and Deployment

Another important area of future work is broader system integration and deployment. The proposed framework could be integrated directly with popular email platforms such as Gmail or Outlook, allowing phishing detection to occur within familiar user environments.

It could also be deployed as a cloud-based service to improve scalability and support large-scale usage. In addition, the development of dedicated APIs would make it easier to integrate the system with other cybersecurity tools, enterprise security platforms, and automated response workflows.

11.5 User Interface Enhancements

Future development may also focus on improving the user interface and overall user experience. A more advanced dashboard could be designed to visualize detection results,

confidence scores, alert categories, and system insights in a clearer and more interactive manner.

Real-time alerts and notification mechanisms could also be added to improve responsiveness for users and analysts. Furthermore, enhancing the explainability features of the system would help users better understand model decisions, thereby increasing trust, usability, and practical effectiveness.

11.6 Expansion to Other Attack Types

The scope of the system can also be extended beyond phishing detection alone. Future versions could be adapted to identify related cyber threats such as spam, malware-distributing links, business email compromise attempts, and broader forms of social engineering.

Integration with external threat intelligence platforms could further strengthen detection capability by providing updated knowledge about malicious domains, known campaigns, and evolving attack behavior. Another promising extension is multilingual phishing detection, which would allow the system to operate across a wider range of linguistic and geographical contexts.

11.7 Summary of Future Scope

The future scope of this project lies in enhancing detection capability, improving scalability, and expanding applicability across real-world cybersecurity environments. By incorporating more advanced models, real-time processing, richer feature engineering, stronger system integration, and broader threat coverage, the proposed framework can evolve into a more comprehensive and intelligent cyber threat detection platform.

Overall, these future directions provide a clear path for transforming the current system from a strong phishing detection solution into a more adaptive, scalable, and versatile security framework for modern digital environments.

REFERENCES

Research Papers and Literature

- [1] *AI Based Phishing Detection Techniques: A Comparative Analysis of Model Performance*. Available: <https://www.researchgate.net/publication/400584536>
- [2] *The Role of Artificial Intelligence in Detecting and Preventing Phishing Emails*. Available: <https://doi.org/10.5281/zenodo.14621440>
- [3] *Hybrid Machine Learning Model for Phishing Detection*. Available: <https://doi.org/10.1109/IS61756.2024.10705257>
- [4] *Phishing Detection System Through Hybrid Machine Learning Based on URL*. Available: <https://doi.org/10.1109/ACCESS.2023.3252366>
- [5] *Across the Spectrum: In-Depth Review of AI-Based Models for Phishing Detection*. Available: <https://doi.org/10.1109/OJCOMS.2024.3462503>
- [6] *AI-Based Phishing Detection and Student Cybersecurity Awareness in the Digital Age*. Available: <https://doi.org/10.1109/OJCOMS.2024.3462503>
- [7] *A Deep Learning-Based Phishing Detection System Using CNN, LSTM, and LSTM-CNN*. Available: <https://doi.org/10.3390/electronics12010232>
- [8] *Hybrid AI Models for Phishing URL Detection: A Unified Approach to Mitigating Cyber Threats*. Available: <https://eudoxuspress.com/index.php/pub/article/view/3026>
- [9] *Neurosymbolic Learning and Domain Knowledge-Driven Explainable AI for Enhanced IoT Network Attack Detection and Response*. Available: <https://doi.org/10.1016/j.cose.2025.104318>

Core Technical References

- [10] V. Metsis, I. Androutsopoulos, and G. Paliouras, “Spam Filtering with Naive Bayes: Which Naive Bayes?,” in *CEAS*, 2006.
- [11] T. Fawcett, “An Introduction to ROC Analysis,” *Pattern Recognition Letters*, 2006.
- [12] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-Training of Deep Bidirectional Transformers for Language Understanding,” in *NAACL*, 2019.
- [13] S. Sanh, L. Debut, J. Chaumond, and T. Wolf, “DistilBERT: A Distilled Version of BERT: Smaller, Faster, Cheaper and Lighter,” *arXiv preprint arXiv:1910.01108*, 2019.
- [14] F. Pedregosa, G. Varoquaux, A. Gramfort, *et al.*, “Scikit-learn: Machine Learning in Python,” *Journal of Machine Learning Research*.

Datasets Used

- [15] J. Nazario, “Phishing Corpus Dataset,” 2005.
- [16] B. Klimt and Y. Yang, “The Enron Corpus: A New Dataset for Email Classification Research,” in *ECML*, 2004.
- [17] *SpamAssassin Public Corpus*, Apache Software Foundation.
- [18] *PhishTank Dataset*. Available: <https://phishtank.org>
- [19] *Top-1M Websites Dataset* (Benign URLs).